

Manipulating Tabular Data

Lesson 5 with *Benoit Parmentier*

Lesson Objectives

- Review what makes a dataset tidy.
- Meet a complete set of functions for most table manipulations.
- Learn to transform datasets with split-apply-combine procedures.
- Understand the basic join operation.

Specific Achievements

- Reshape data frames with pandas
- Summarize data by groups with pandas
- Combine multiple data frame operations with method chaining (piping with pandas “.”)
- Combine multiple data frames with “joins” (merge)

Data frames occupy a central place in Python data analysis pipelines. The pandas package provides the objects and most necessary tools to subset, reformat and transform data frames. The key functions in the package have close counterparts in SQL (Structured Query Language), which provides the added bonus of facilitating translation between python and relational databases.

Tidy Concept

Most time is spent on cleaning and wrangling data rather than analysis. In 2014, Hadley Wickam (R developer at RStudio) published a paper that defines the concepts underlying tidy datasets. Hadley Wickam defined tidy datasets as those where:

- each variable forms a column (also called field)
- each observation forms a row
- each type of observational unit forms a table

These guidelines may be familiar to some of you—they closely map to best practices for “normalization” in database design. It corresponds to the 3rd normal form’s described by Codd 1990 but uses the language of statistical analysis rather than relational database.

Consider a data set where the outcome of an experiment has been recorded in a perfectly appropriate way:

bloc	drug	control	placebo
1	0.22	0.58	0.31
2	0.12	0.98	0.47
3	0.42	0.19	0.40

The response data are present in a compact matrix, as you might record it on a spreadsheet. The form does not match how we think about a statistical model, such as:

$$response \sim block + treatment$$

In a tidy format, each row is a complete observation: it includes the response value and all the predictor values. In this data, some of those predictor values are column headers, so the table needs to be reshaped. The pandas package provides functions to help re-organize tables.

The third principle of tidy data, one table per category of observed entities, becomes especially important in synthesis research. Following this principle requires holding tidy data in multiple tables, with associations between them formalized in metadata, as in a relational database.

Datasets split across multiple tables are unavoidable in synthesis research, and commonly used in the following two ways (often in combination):

- two tables are “un-tidied” by joins, or merging them into one table
- statistical models conform to the data model through a hierarchical structure or employing “random effects”

The pandas package includes several functions that all perform variations on table joins needed to “un-tidy” your tables, but there are only two basic types of table relationships to recognize:

- **One-to-one** relationships allow tables to be combined based on the same unique identifier (or “primary key”) in both tables.
- **Many-to-one** relationships require non-unique “foreign keys” in the first table to match the primary key of the second.

[Top of Section](#)

Wide to Long

The `pandas` package’s `melt` function reshapes “wide” data frames into “long” ones.

```
worksheet-5.ipynb
```

```
import pandas as pd
import numpy as np
trial_df = pd.DataFrame({"block": [1,2,3],
                          "drug": [0.22,0.12,0.42],
                          "control": [0.58,0.98,0.19],
                          "placebo": [0.31,0.47,0.40]})
trial_df.head()
```

	block	control	drug	placebo
0	1	0.58	0.22	0.31
1	2	0.98	0.12	0.47
2	3	0.19	0.42	0.40

```
worksheet-5.ipynb
```

```
tidy_trial_df = pd.melt(trial_df,
                        id_vars=['block'],
                        var_name='treatment',
                        value_name='response')
tidy_trial_df.head()
```

	block	treatment	response
0	1	control	0.58
1	2	control	0.98
2	3	control	0.19
3	1	drug	0.22
4	2	drug	0.12

All columns, except for “block”, are stacked in two columns: a “key” and a “value”. The key column gets the name `treatment` and the value column receives the name `response`. For each row in the result, the key is taken from the name of the column and the value from the data in the column.

Long to Wide

Data can also fail to be tidy when a table is too long. The Entity-Attribute-Value (EAV) structure common in large databases distributes multiple attributes of a single entity/observation into separate rows.

Remember that the exact state of “tidy” may depend on the analysis: the key is knowing what counts as a complete observation. For example, the community ecology package `vegan` requires a matrix of species counts, where rows correspond to species and columns to sites. This may seem like too “wide” a format, but in the packages several multi-variate analyses, the abundance of a species across multiple sites is considered a complete observation.

Consider survey data on participant’s age and income stored in a EAV structure.

```
worksheet-5.ipynb
```

```
df2 = tidy_trial_df.pivot(index='block',
                          columns='treatment',
                          values='response')
```

```
df2 = df2.reset_index()
df2.columns
```

```
Index(['block', 'control', 'drug', 'placebo'], dtype='object', name='treatment')
```

worksheet-5.ipynb

```
df2.reset_index()
```

	treatment	index	block	control	drug	placebo
0		0	1	0.58	0.22	0.31
1		1	2	0.98	0.12	0.47
2		2	3	0.19	0.42	0.40

worksheet-5.ipynb

```
df2
```

	treatment	block	control	drug	placebo
0		1	0.58	0.22	0.31
1		2	0.98	0.12	0.47
2		3	0.19	0.42	0.40

Consider survey data on participant's age and income *stored* in a EAV structure.

worksheet-5.ipynb

```
from io import StringIO, BytesIO
```

```
text_string = StringIO("""
participant,attr,val
1,age,24
2,age,57
3,age,13
1,income,30
2,income,60
""")
```

```
survey_df = pd.read_csv(text_string, sep=",")
survey_df
```

	participant	attr	val
0	1	age	24
1	2	age	57
2	3	age	13
3	1	income	30
4	2	income	60

Transform the data with the `pivot` function, which “reverses” a `melt`. These are equivalent to `spread` and `gather` in the dplyr r package.

worksheet-5.ipynb

```
tidy_survey = survey_df.pivot(index='participant',
                               columns='attr',
                               values='val')
print(tidy_survey.head())
```

	attr	age	income
participant			
1		24.0	30.0
2		57.0	60.0
3		13.0	NaN

worksheet-5.ipynb

tidy_survey = tidy_survey.reset_index()
tidy_survey.columns

Index(['participant', 'age', 'income'], dtype='object', name='attr')

worksheet-5.ipynb

tidy_survey.reset_index()

attr	index	participant	age	income
0	0	1	24.0	30.0
1	1	2	57.0	60.0
2	2	3	13.0	NaN

worksheet-5.ipynb

tidy_survey

attr	participant	age	income
0	1	24.0	30.0
1	2	57.0	60.0
2	3	13.0	NaN

[Top of Section](#)

Sample Data



Credit: US Census Bureau

To learn about data transformation with pandas, we need more data. The Census Bureau collects subnational economic data for the U.S., releasing annual [County Business Patterns \(CBP\)](#) datasets including the number of establishments, employment, and payroll by industry. They also conduct the [American Community Survey \(ACS\)](#) and publish, among other demographic and economic variables, estimates of median income for individuals working in different industries.

- [County Business Patterns \(CBP\)](#)
- [American Community Survey \(ACS\)](#)

worksheet-5.ipynb

import pandas as pd
cbp = pd.read_csv('data/cbp15co.csv')
cbp.describe()

worksheet-5.ipynb

print(cbp.dtypes)

FIPSTATE int64
FIPSCITY int64

NAICS	object
EMPFLAG	object
EMP_NF	object
EMP	int64
QP1_NF	object
QP1	int64
AP_NF	object
AP	int64
EST	int64
N1_4	int64

See the [CBP dataset documentation](#) for an explanation of the variables we don't discuss in this lesson.

Modify the import to clean up this read: consider the data type for FIPS codes along with what string in this CSV file represents NAs, a.k.a. data that is not-available or missing.

```
worksheet-5.ipynb

import numpy as np
import pandas as pd

cbp = pd.read_csv(
    'data/cbp15co.csv',
    na_values = "NULL",
    keep_default_na=False,
    dtype = {"FIPSTATE": np.str,
            "FIPSCTY": np.str}
)
```

Question

What changed?

Answer

Using `dtypes()` shows that the character string "" in the CSV file is no longer read into R as missing data (an NA) but as an empty string. The two named "FIPS" columns are now correctly read as strings.

```
worksheet-5.ipynb

import pandas as pd
import numpy as np
acs = pd.read_csv(
    'data/ACS/sector_ACS_15_5YR_S2413.csv',
    dtype = {"FIPS": np.str}
)
```

Now let's display the data types

```
worksheet-5.ipynb

#acs.dtypes
print(acs.dtypes)
```

```
FIPS          object
County        object
Sector        object
median_income float64
dtype: object
```

The two datasets both contain economic variables for each U.S. county and specified by different categories of industry. The data could potentially be manipulated into a single table reflecting the follow statistical model.

$$\text{median_income} \sim \text{industry} + \text{establishment_size}$$



Key Functions

Function	Returns
<code>query</code>	keep rows that satisfy conditions
<code>assign</code>	apply a transformation to existing [split] columns
<code>['col1', 'col2']</code>	select and keep columns with matching names
<code>merge</code>	merge columns from separate tables into one table
<code>groupby</code>	split data into groups by an existing factor
<code>agg</code>	summarize across rows to use after groupby [and combine split groups]

The table above summarizes the most commonly used functions in [pandas](#), which we will demonstrate in turn on data from the U.S. Census Bureau.

Filter and pattern matching

The `cbp` table includes character `NAICS` column. Of the 2 million observations, lets see how many observations are left when we keep only the 2-digit NAICS codes, representing high-level sectors of the economy.

worksheet-5.ipynb  



```
#import pandas as pd
cbp2 = cbp[cbp['NAICS'].str.contains("----")]
cbp2 = cbp2[~cbp2.NAICS.str.contains("----")]
cbp2.head()
```

	FIPSTATE	FIPSCTY	NAICS	EMPFLAG	...	N1000_3	N1000_4	CENSTATE	CENCTY
1	01	001	11----		...	0	0	63	1
10	01	001	21----		...	0	0	63	1
17	01	001	22----		...	0	0	63	1
27	01	001	23----		...	0	0	63	1
93	01	001	31----		...	0	0	63	1

[5 rows x 26 columns]

Note that a logical we used the function `contains` from pandas to filter the dataset in two steps. The function contains allows for pattern matching of any character within strings. The `~` is used to remove the rows that contains specific patterns.

Filtering string often uses pattern matching by [regular expressions](#) which may be a bit more manageable, and streamlines the operations.

worksheet-5.ipynb  

```
cbp3 = cbp[cbp['NAICS'].str.contains('[0-9]{2}----')]
cbp3.head()
```

	FIPSTATE	FIPSCTY	NAICS	EMPFLAG	...	N1000_3	N1000_4	CENSTATE	CENCTY
1	01	001	11----		...	0	0	63	1
10	01	001	21----		...	0	0	63	1
17	01	001	22----		...	0	0	63	1
27	01	001	23----		...	0	0	63	1
93	01	001	31----		...	0	0	63	1

[5 rows x 26 columns]

Altering, updating and transforming columns

The `assign` function is the [pandas](#) answer to updating or altering your columns. It performs arbitrary operations on existing columns and appends the result as a new column of the same length.

Here are two ways to create a new column using `assign` and the `[]` operators.

worksheet-5.ipynb



```
cbp3["FIPS"] = cbp3["FIPSTATE"]+cbp3["FIPSCTY"]
```

```
/usr/bin/python3:1: SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.
```

```
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing
```

worksheet-5.ipynb



```
cbp3.assign(FIPS2=lambda x: x['FIPSTATE']+x['FIPSCTY'])
```

	FIPSTATE	FIPSCTY	NAICS	EMPFLAG	...	CENSTATE	CENCTY	FIPS	FIPS2
1	01	001	11----		...	63	1	01001	01001
10	01	001	21----		...	63	1	01001	01001
17	01	001	22----		...	63	1	01001	01001
27	01	001	23----		...	63	1	01001	01001
93	01	001	31----		...	63	1	01001	01001
163	01	001	42----		...	63	1	01001	01001
218	01	001	44----		...	63	1	01001	01001
351	01	001	48----		...	63	1	01001	01001
381	01	001	51----		...	63	1	01001	01001
401	01	001	52----		...	63	1	01001	01001
429	01	001	53----		...	63	1	01001	01001
465	01	001	54----		...	63	1	01001	01001

worksheet-5.ipynb



```
cbp3.shape
```

```
(58901, 27)
```

worksheet-5.ipynb



```
cbp3.head()
```

	FIPSTATE	FIPSCTY	NAICS	EMPFLAG	...	N1000_4	CENSTATE	CENCTY	FIPS
1	01	001	11----		...	0	63	1	01001
10	01	001	21----		...	0	63	1	01001
17	01	001	22----		...	0	63	1	01001
27	01	001	23----		...	0	63	1	01001
93	01	001	31----		...	0	63	1	01001

```
[5 rows x 27 columns]
```

Select

To keep particular columns of a data frame (rather than filtering rows), use the `filter` or `[]` functions with arguments that match column names.

worksheet-5.ipynb



```
cbp2.columns
```

```
Index(['FIPSTATE', 'FIPSCTY', 'NAICS', 'EMPFLAG', 'EMP_NF', 'EMP', 'QP1_NF',  
      'QP1', 'AP_NF', 'AP', 'EST', 'N1_4', 'N5_9', 'N10_19', 'N20_49',  
      'N50_99', 'N100_249', 'N250_499', 'N500_999', 'N1000', 'N1000_1',  
      'N1000_2', 'N1000_3', 'N1000_4', 'CENSTATE', 'CENCTY'],  
      dtype='object')
```

One way to “match” is by including complete names, each one you want to keep:

worksheet-5.ipynb



```
cbp3 = cbp3[['FIPS', 'NAICS', 'N1_4', 'N5_9', 'N10_19']]
cbp3.head()
```

	FIPS	NAICS	N1_4	N5_9	N10_19
1	01001	11----	5	1	0
10	01001	21----	0	1	1
17	01001	22----	2	1	2
27	01001	23----	51	13	7
93	01001	31----	9	4	4

Alternatively, we can use the `filter` function to select all columns starting with N or matching with 'FIPS' or 'NAICS' pattern. The `filter` command is useful when chaining methods (or piping operations).

worksheet-5.ipynb

```
cbp4= cbp.filter(regex='^N|FIPS|NAICS',axis=1)
cbp4.head()
```

	FIPSTATE	FIPSCTY	NAICS	N1_4	...	N1000_1	N1000_2	N1000_3	N1000_4
0	01	001	-----	430	...	0	0	0	0
1	01	001	11----	5	...	0	0	0	0
2	01	001	113///	4	...	0	0	0	0
3	01	001	1133//	4	...	0	0	0	0
4	01	001	11331/	4	...	0	0	0	0

[5 rows x 16 columns]

[Top of Section](#)

Join

The CBP dataset uses FIPS to identify U.S. counties and NAICS codes to identify types of industry. The ACS dataset also uses FIPS but their data may aggregate across multiple NAICS codes representing a single industry sector.

worksheet-5.ipynb

```
sector = pd.read_csv(
    'data/ACS/sector_naics.csv',
    dtype = {"NAICS": np.int64})
print(sector.dtypes)
```

```
Sector    object
NAICS      int64
dtype: object
```

worksheet-5.ipynb

```
print(cbp.dtypes)
```

```
FIPSTATE    object
FIPSCTY     object
NAICS       object
EMPFLAG     object
EMP_NF      object
EMP         int64
QP1_NF      object
QP1         int64
AP_NF       object
AP          int64
EST         int64
N1_4        int64
```


worksheet-5.ipynb



```
cbp.head()
```

	FIPSTATE	FIPSCTY	NAICS	EMPFLAG	...	N1000_3	N1000_4	CENSTATE	CENCTY
0	01	001	-----		...	0	0	63	1
1	01	001	11----		...	0	0	63	1
2	01	001	113///		...	0	0	63	1
3	01	001	1133//		...	0	0	63	1
4	01	001	11331/		...	0	0	63	1

[5 rows x 26 columns]

worksheet-5.ipynb



```
cbp.dtypes
```

```
FIPSTATE    object
FIPSCTY     object
NAICS       object
EMPFLAG     object
EMP_NF      object
EMP         int64
QP1_NF      object
QP1         int64
AP_NF       object
AP          int64
EST         int64
N1_4        int64
N5_9        int64
```

worksheet-5.ipynb



```
cbp.head()
```

	FIPSTATE	FIPSCTY	NAICS	EMPFLAG	...	N1000_3	N1000_4	CENSTATE	CENCTY
0	01	001	-----		...	0	0	63	1
1	01	001	11----		...	0	0	63	1
2	01	001	113///		...	0	0	63	1
3	01	001	1133//		...	0	0	63	1
4	01	001	11331/		...	0	0	63	1

[5 rows x 26 columns]

worksheet-5.ipynb



```
print(sector.dtypes)
```

```
Sector      object
NAICS       int64
dtype: object
```

worksheet-5.ipynb



```
print(sector.shape) #24 economic sectors
```

(24, 2)

worksheet-5.ipynb



```
sector.head()
```

	Sector	NAICS
0	agriculture forestry fishing and hunting	11
1	mining quarrying and oil and gas extraction	21

2	utilities	22
3	construction	23
4	manufacturing	31

Probably the primary challenge in combining secondary datasets for synthesis research is dealing with their different sampling frames. A very common issue is that data are collected at different “scales”, with one dataset being at higher spatial or temporal resolution than another. The differences between the CBP and ACS categories of industry present a similar problem, and require the same solution of re-aggregating data at the “lower resolution”.

Many-to-One

Before performing the join operation, some preprocessing is necessary to extract from the NAICS columns the first two digits matching the sector identifiers.

worksheet-5.ipynb

```
logical_idx = cbp['NAICS'].str.match('[0-9]{2}----') #boolean index
cbp = cbp.loc[logical_idx]
cbp.head()
```

	FIPSTATE	FIPSCTY	NAICS	EMPFLAG	...	N1000_3	N1000_4	CENSTATE	CENCTY
1	01	001	11----		...	0	0	63	1
10	01	001	21----		...	0	0	63	1
17	01	001	22----		...	0	0	63	1
27	01	001	23----		...	0	0	63	1
93	01	001	31----		...	0	0	63	1

[5 rows x 26 columns]

worksheet-5.ipynb

```
cbp.shape
```

(58901, 26)

worksheet-5.ipynb

```
cbp['NAICS'] = cbp.NAICS.apply(lambda x: np.int64(x[0:2])) # select first two digits
```

worksheet-5.ipynb

```
#Many to one to join economic sector code to NAICS
```

```
cbp_test = cbp.merge(sector, on = "NAICS", how='inner')
cbp_test.head()
```


	FIPSTATE	FIPSCTY	...	CENCTY	Sector
0	01	001	...	1	agriculture forestry fishing and hunting
1	01	003	...	3	agriculture forestry fishing and hunting
2	01	005	...	5	agriculture forestry fishing and hunting
3	01	007	...	7	agriculture forestry fishing and hunting
4	01	009	...	9	agriculture forestry fishing and hunting

[5 rows x 27 columns]

worksheet-5.ipynb

```
print(cbp_test.shape)
```

(56704, 27)



The NAICS field in the `cbp` table can have the same value multiple times, it is not a primary key in this table. In the `sector` table, the NAICS field is the primary key uniquely identifying each record. The type of relationship between these tables is therefore “many-to-one”.

Question

Note that we lost a couple thousand rows through this join. How could `cbp` have fewer rows after a join on NAICS codes?

Answer

The CBP data contains an NAICS code not mapped to a sector—the “error code” 99 is not present in `sector`. The use of “error codes” that could easily be mistaken for data is frowned upon.

Group By

A very common data manipulation procedure known as “split-apply-combine” tackles the problem of applying the same transformation to subsets of data while keeping the result all together. We need the total number of establishments in each size class *aggregated within* each county and industry sector.

The pandas function `groupby` begins the process by indicating how the data frame should be split into subsets.

```

worksheet-5.ipynb

cbp["FIPS"] = cbp["FIPSTATE"]+cbp["FIPSCTY"]
cbp = cbp.merge(sector, on = "NAICS")

cbp_grouped = cbp.groupby(['FIPS', 'Sector'])
cbp_grouped

<pandas.core.groupby.generic.DataFrameGroupBy object at 0x7f5974d485f8>
  
```

At this point, nothing has really changed:

```

worksheet-5.ipynb

cbp_grouped.dtypes

FIPS  Sector  FIPSTATE  ...  CENCTY
01001  accommodation and food services  object  ...  int64
      administrative and support and waste management...  object  ...  int64
      agriculture forestry fishing and hunting  object  ...  int64
      arts entertainment and recreation  object  ...  int64
      construction  object  ...  int64
      educational services  object  ...  int64
      finance and insurance  object  ...  int64
      health care and social assistance  object  ...  int64
  
```

information
management of companies and enterprises

object ... int64
object ... int64

The `groupby` statement generates a groupby data frame. You can add multiple variables (separated by commas) in `groupby`; each distinct combination of values across these columns defines a different group.

Summarize

The operation to perform on each group is summing: we need to sum the number of establishments in each group. Using `pandas` functions, the summaries are automatically combined into a data frame.

worksheet-5.ipynb


```
grouped_df = (cbp
.groupby(['FIPS', 'Sector'])
.agg('sum')
.filter(regex='^N')
.drop(columns=['NAICS'])
)
```

```
grouped_df.head(5)
```

FIPS	Sector	N1_4	...	N1000_4
01001	accommodation and food services	23	...	0
	administrative and support and waste management...	18	...	0
	agriculture forestry fishing and hunting	5	...	0
	arts entertainment and recreation	5	...	0
	construction	51	...	0

[5 rows x 13 columns]

The “combine” part of “split-apply-combine” occurs automatically, when the attributes introduced by `groupby` are dropped. You can see attributes by running the `dtypes` function on the data frame.



There is now a one-to-one relationship between `cbp` and `acs`, based on the combination of FIPS and Sector as the primary key for both tables.

worksheet-5.ipynb

```
print(grouped_df.shape)
```

(56704, 13)

worksheet-5.ipynb



```
print(acs.shape)
```

```
(59698, 4)
```

worksheet-5.ipynb



```
acs_cbp = grouped_df.merge(acs,on='FIPS',)  
print(acs_cbp.shape)
```

```
(1061416, 17)
```

Again, however, the one-to-one relationship does not mean all rows are preserved by the join. The specific nature of the `inner_join` is to keep all rows, even duplicating rows if the relationship is many-to-one, where there are matching values in both tables, and discarding the rest.

The `acs_cbp` table now includes the `median_income` variable from the ACS and appropriately aggregated establishment size information (the number of establishments by employee bins) from the CBP table.

worksheet-5.ipynb



```
acs_cbp.head()
```

	FIPS	N1_4	...	Sector	median_income
0	01001	23	...	agriculture forestry fishing and hunting	27235.0
1	01001	23	...	mining quarrying and oil and gas extraction	102722.0
2	01001	23	...	construction	31632.0
3	01001	23	...	manufacturing	40233.0
4	01001	23	...	wholesale trade	41656.0

```
[5 rows x 17 columns]
```

[Top of Section](#)




Additional Resources

The following cheat sheets and tutorials repeat much of this lesson, but also provide information on additional functions for “data wrangling”.

- [Data Wrangling Cheat Sheet](#)
- [Tidyverse In Pandas](#)
- [String and Text With Pandas](#)

The first is a set of cheat sheets created by [pydata.org](#), and provides a handy, visual summary of all the key functions discussed in this lesson. It also lists some of the auxiliary functions that can be used within each type of expression, e.g. aggregation functions for summarize, “moving window” functions for mutate, etc. For those familiar with the tidyverse univers, please consult the second link.

[Top of Section](#)

If you need to catch-up before a section of code will work, just squish it's  to copy code above it into your clipboard. Then paste into your interpreter's console, run, and you'll be ready to start in on that section. Code copied by both  and  will also appear below, where you can edit first, and then copy, paste, and run again.

```
# Nothing here yet!
```